

# Description of Polgraw All-Sky Search GPGPU program

Jan Bolek

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	The algorithm . . . . .	2
<b>2</b>	<b>Code explanation</b>	<b>3</b>
2.1	Data structures . . . . .	3
2.2	"main" function . . . . .	3
2.3	Initialization . . . . .	4
2.4	Main job . . . . .	4
2.5	Sky loop internals . . . . .	5
2.5.1	Preparing variables . . . . .	5
2.5.2	Resampling and spline interpolation to baricentric time . . . . .	6
2.6	Spin-down loop internals . . . . .	8
2.7	Finding and saving candidates data . . . . .	10
<b>3</b>	<b>References</b>	<b>12</b>

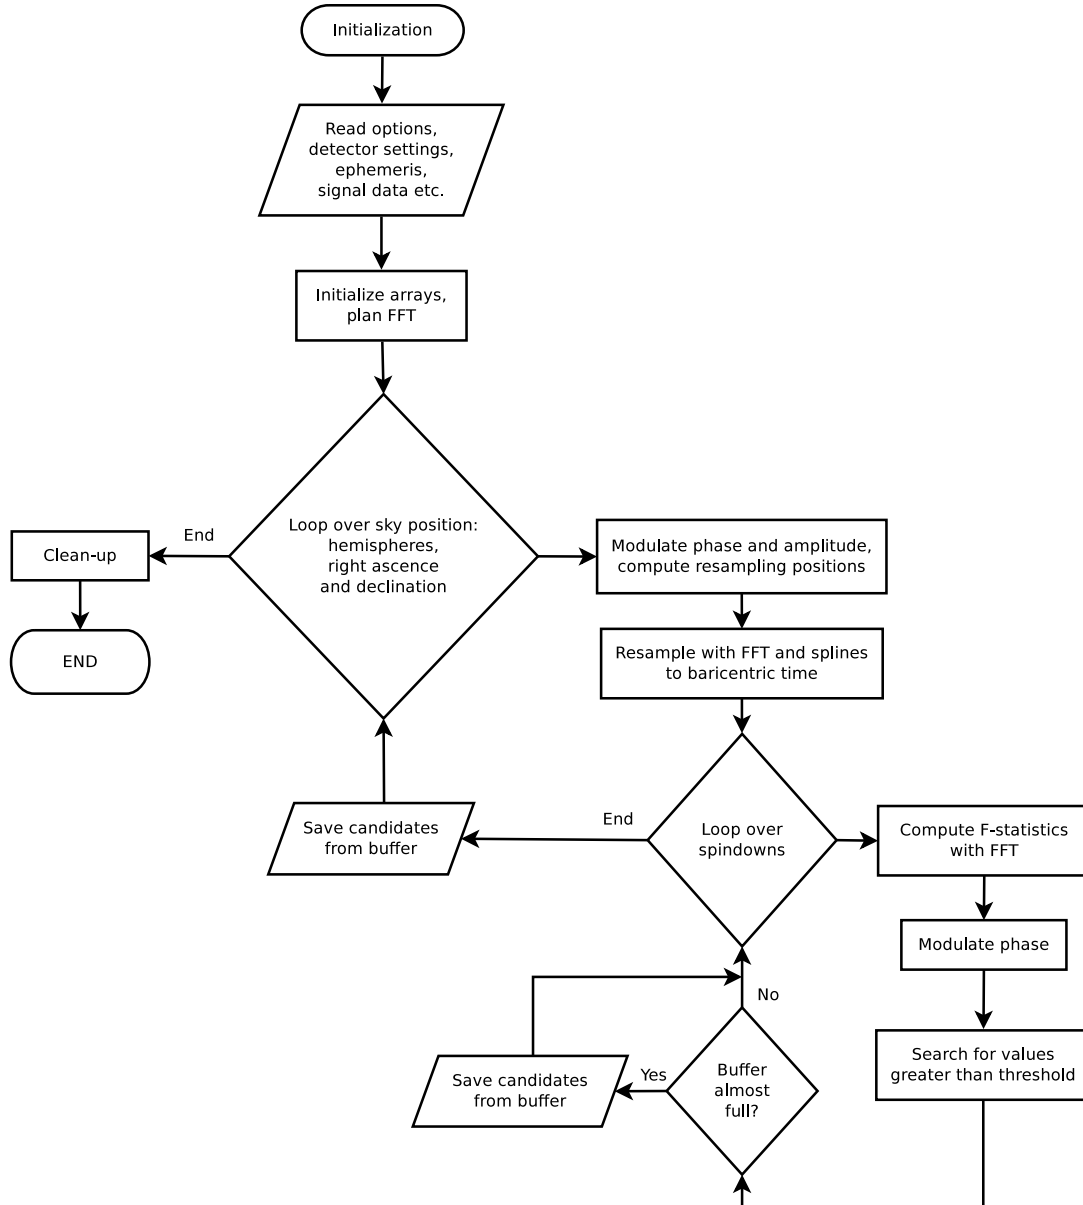
# 1 Introduction

The purpose of this article is to introduce one into structure, working principle and techniques used in gravitational waves (abbr. GW) All-Sky Search program.

The main purpose of described program is to find candidate GW signals parameters using  $\mathcal{F}$ -statistic. The parameters consist of frequency, spin-down (first derivative of frequency), positions on sky (e.g. right ascension and declination) and signal-to-noise ratio (SNR). Further theoretical explanation can be found in [1].

## 1.1 The algorithm

Figure 1: Flow-chart of the algorithm



The searching algorithm consists of iterating over sky positions (i.e. hemispheres, right ascension and declination transformed into integral-valued grid), frequency and one spin-down parameter in order to modulate signal with appropriate function using this parameters and evaluate so-called  $\mathcal{F}$ -statistic. The trick is that  $\mathcal{F}$ -statistic may be computed using Fast Fourier Transform (FFT), thus iterating over frequency loop can be avoided, which results in great performance improvement.

In order to apply above method detector signal must be defined over barycentric time (i.e. time with respect to solar system barycenter), so there must be performed interpolation to precomputed barycentric time positions. Those positions are known from detector's and Earth's ephemeris.

Because of usage of large vectors ( $\sim 2^{19}$  elements) and Fourier transform the program was written for GPUs using CUDA architecture. Thanks to use of massive parallelization there was achieved speed-up over CPU version of about 50-100 times, depending on GPU type.

## 2 Code explanation

### 2.1 Data structures

For convenience arrays and other variables, e.g. settings and detector constants are stored in structures defined in file `struct.h`. Basic structures are:

- `struct _comm_line_opts` (typedef `Command_line_opts`)  
options read from command-line
- `struct _detector_settings` (typedef `Detector_settings`)  
parameters and settings for detector and for computations (array lengths etc.)
- `struct _search_range` (typedef `Search_range`)  
range of search
- `struct _fft_plans` (typedef `FFT_plans`)  
plans (handles) for cuFFT
- `struct _ampl_mod_coeff` (typedef `Ampl_mod_coeff`)  
amplitude modulation function coefficients

### 2.2 "main" function

`main()` function consists of other functions calls – mostly initialization. The most important part is:

Listing 1: `main.c`

```
1  /*
2  ##### Main job #####
```

```

3  */
4
5  search(
6      &sett,
7      &opts,
8      &s_range,
9      &arr,
10     &fft_plans,
11     &amod,
12     &Fnum,
13     cu_F
14 );

```

i.e. the worker function call.

## 2.3 Initialization

Initialization consists of the following steps:

1. handling command line options (`handle_opts()`)
2. checking if output directory exists (and creating it if needed)
3. reading grid generating matrix (`read_grid()`)
4. reading detector settings and physical constants (`settings()` from `settings.c`)
5. initializing amplitude modulation function coefficients for Virgo detector (`rogcvir()`)
6. allocating arrays and reading input arrays (`init_arrays()`)
7. determining searching range (`set_search_range()`)
8. planning FFT (`plan_fft()`)
9. reading checkpoints (for resuming execution at given point) (`read_checkpoints()`)

All those functions, except for `settings()`, are defined in `init.c`.

## 2.4 Main job

Function `search()` is defined in file `jobcore.cu`. It contains main loop over sky positions and `jobcore()` function call.

`jobcore()` function performs computations outside of spin-down loop and in the spin-down loop itself.

## 2.5 Sky loop internals

### 2.5.1 Preparing variables

First, we compute real angular coordinates from linear grid coordinates:

Listing 2: jobcore.cu

```
202 //change linear (grid) coordinates to real coordinates
203 lin2ast (a11/sett->oms, a12/sett->oms, pm, sett->sepsm, sett->cepsm,
204         &sinlt, &cosalt, &sindelt, &cosdelt);
205
206 // calculate declination and right ascencion
207 // which will be written in file as candidate signal sky positions
208 sgnlt[2] = asin (sindelt);
209 sgnlt[3] = fmod (atan2 (sinalt, cosalt)+2.*M_PI, 2.*M_PI);
```

Then we compute amplitude modulation functions  $a(t)$  and  $b(t)$  (`arr->cu_aa` and `arr->cu_bb`, respectively) using GPU. In this step we need to normalize those function by square root of their sum of squares (i.e.  $a_i = a'_i / \sqrt{\sum_t a_i'^2}$ ), which is performed using parallel reduction in kernels `reduction_sumsq()` and `reduction_sum()`.

Listing 3: jobcore.cu

```
212 // amplitude modulation function
213 modvir_gpu(sinalt, cosalt, sindelt, cosdelt, sett->sphir,
214           sett->cphir, arr->cu_aa, arr->cu_bb, sett->N, arr);
```

Then kernel `shift_time_mod()` computes shifted (baricentric) time, modulates signal with  $a(t)$ ,  $b(t)$  and phase factor (thus obtaining two vectors, `xa` and `xb`), computes shifted time position `tshift[i]` for spline interpolation and pads zeros to nearest power of two (`nfft`).

Listing 4: jobcore.cu

```
216 //calculate detector positions with respect to baricenter
217 nSource[0] = cosalt*cosdelt;
218 nSource[1] = sinalt*cosdelt;
219 nSource[2] = sindelt;
220
221 shft1 = 0.;
222 for (j=0; j<3; j++)
223     shft1 += nSource[j]*arr->DetSSB[j];
224 het0 = fmod (nn*sett->M[8]+mm*sett->M[12], sett->M[0]);
225
226 //compute shifted time, modulate phase and pad zeros to size 'nfft'
227 shift_time_mod<<<(sett->nfft + BLOCK_SIZE - 1)/BLOCK_SIZE, BLOCK_SIZE
228 >>>(shft1, het0, nSource[0], nSource[1], nSource[2],
229     arr->cu_xDat, arr->cu_xa, arr->cu_xb, arr->cu_shft,
230     arr->cu_shftf, arr->cu_tshift, arr->cu_aa, arr->cu_bb, arr->
231     cu_DetSSB,
232     sett->oms, sett->N, sett->nfft, sett->interpftpad);
```

Listing 5: kernels.cu

```
42 __global__ void shift_time_mod(...) {
43     int i = blockIdx.x * blockDim.x + threadIdx.x;
```

```

44  if (i < N) {
45      double S = ns0 * DetSSB[i*3+0] + ns1 * DetSSB[i*3+1] + ns2 *
          DetSSB[i*3+2];
46      shft[i] = S;
47      shftf[i] = S - shft1;
48
49      /* phase mod */
50      double phase = -het0*i - oms * S;
51      double c = cos(phase), s = sin(phase);
52      xa[i].x = xDat[i] * aa[i] * c;
53      xa[i].y = xDat[i] * aa[i] * s;
54      xb[i].x = xDat[i] * bb[i] * c;
55      xb[i].y = xDat[i] * bb[i] * s;
56
57      //calculate time positions for spline interpolation
58      tshift[i] = interpftpad * ( i - shftf[i] );
59  } else if (i < nfft) {
60      xa[i].x = xa[i].y = xb[i].x = xb[i].y = 0;
61  }
62  }

```

### 2.5.2 Resampling and spline interpolation to baricentric time

First we compute forward FFT of length `nfft` (with zeros padded from original length `N`). Then we pad Fourier transform with zeros to size  $2*nfft$  (remembering that in unshifted FFT positive frequencies are placed from 0th bin to  $n/2-1$ , and negative are placed from  $n/2$  to  $n-1$ , so we basically insert zeros "in-between") and take Fourier transform back to obtain signal samples at twice the original sampling frequency (see figure 2.) Then the vector is scaled because of property of FFT (i.e. elements after two FFTs are multiplied by size of vector).

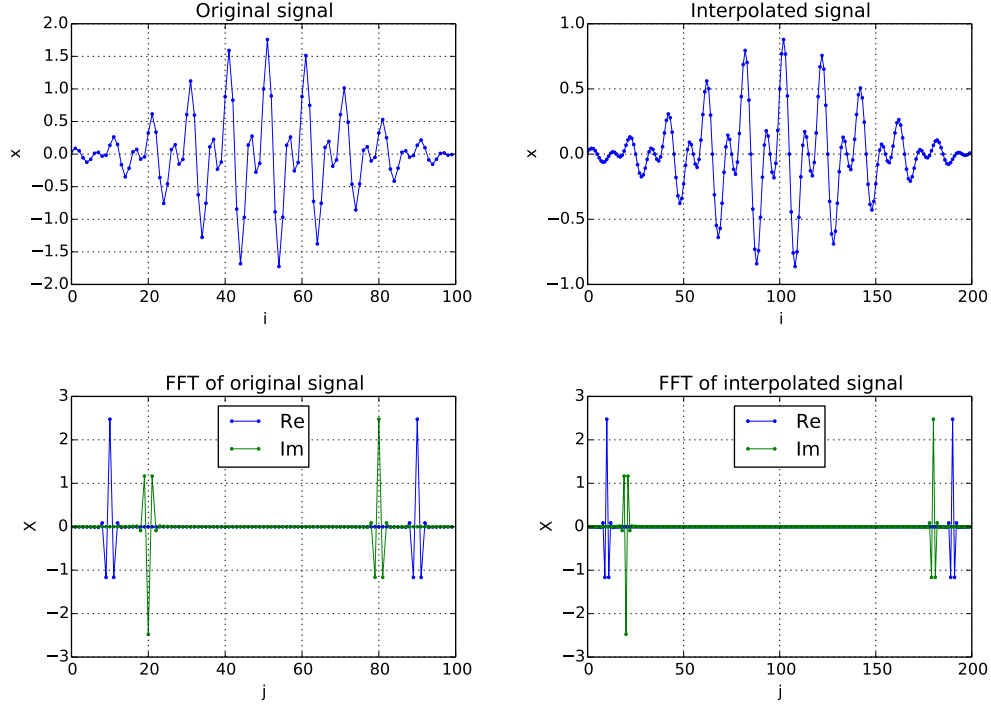
Listing 6: jobcore.cu

```

240  int nyqst = (sett->nfft+2)/2; // Nyquist frequency
241
242  // fft
243  cufftExecZ2Z(plans->pl_int, arr->cu_xa, arr->cu_xa, CUFFT_FORWARD);
244  cufftExecZ2Z(plans->pl_int, arr->cu_xb, arr->cu_xb, CUFFT_FORWARD);
245
246  //shift frequencies and remove those over Nyquist
247  resample_postfft<<<(sett->Ninterp + BLOCK_SIZE - 1)/BLOCK_SIZE,
          BLOCK_SIZE>>>(arr->cu_xa, arr->cu_xb, sett->nfft, sett->Ninterp,
          nyqst);
248  CudaCheckError();
249
250  // fft back
251  cufftExecZ2Z(plans->pl_inv, arr->cu_xa, arr->cu_xa, CUFFT_INVERSE);
252  cufftExecZ2Z(plans->pl_inv, arr->cu_xb, arr->cu_xb, CUFFT_INVERSE);
253
254  //scale fft
255  ft = (double)sett->interpftpad / sett->Ninterp; //scale FFT
256  scale_fft<<<(sett->Ninterp + BLOCK_SIZE - 1)/BLOCK_SIZE, BLOCK_SIZE
          >>>(arr->cu_xa, arr->cu_xb, ft, sett->Ninterp);
257  CudaCheckError();

```

Figure 2: FFT interpolation example



Next step is spline interpolation in which we use cubic splines with natural boundary conditions, meaning that second derivatives at boundaries ( $x_0$  and  $x_n$ ) are set to zero. In our case we also have other assumption:  $h = x_{i+1} - x_i = 1$ .

Spline coefficients are computed by solving linear equation:

$$\begin{bmatrix} 4 & 1 & & & & \\ 1 & 4 & 1 & & & \\ & 1 & \ddots & \ddots & & \\ & & \ddots & \ddots & 1 & \\ & & & 1 & 4 & 1 \\ & & & & 1 & 4 \end{bmatrix} \begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ \vdots \\ z_{n-3} \\ z_{n-2} \\ z_{n-1} \end{bmatrix} = \begin{bmatrix} 6(b_1 - b_0) \\ 6(b_2 - b_1) \\ 6(b_3 - b_2) \\ \vdots \\ 6(b_{n-3} - b_{n-4}) \\ 6(b_{n-2} - b_{n-3}) \\ 6(b_{n-1} - b_{n-2}) \end{bmatrix}$$

where:

- $n$  is number of intervals in interpolated function, so  $n + 1$  is number of given points (knots)
- $z_i$  are unknown coefficients (second derivatives at knots, to be precise),
- $n - 2$  is rank of set of equations (we know from natural boundary conditions that  $z_0 = z_n = 0$ )
- $b_i$  are differences of function values in consecutive knots:  $b_i = (y_{i+1} - y_i)$

Matrix of above shape is called *tridiagonal matrix*. We solve this set of equations using `cusparseZgtsv` function included in the `cuSPARSE` package.

Having solved above equations for  $z_i$ , we can compute interpolated value  $y'(x)$  using following formula:

$$y(x) = S_i(x) = \frac{z_{i+1}}{6}(x-x_i)^3 + \frac{z_i}{6}(x_{i+1}-x)^3 + \left(y_{i+1} - \frac{z_{i+1}}{6}\right)(x-x_i) + \left(y_i - \frac{z_i}{6}\right)(x_{i+1}-x)$$

where  $i$  is number of spline interval  $S_i$  which contains point  $x$ :  $i = \lfloor x \rfloor$ .

Function `gpu_interp()` encapsulates above algorithm. It is defined in file `spline_z.cu`.

Listing 7: `jobcore.cu`

```

261 //interpolation of xa
262 gpu_interp(arr->cu_xa, //input data
263           sett->Ninterp, //input data length
264           arr->cu_tshift, //output time domain
265           arr->cu_xar, //output values
266           sett->N, //output data length
267           arr->cu_d, //diagonal
268           arr->cu_dl, //upper diagonal
269           arr->cu_du, //lower diagonal
270           arr->cu_B); //coefficient matrix
271
272 //interpolation of xb
273 gpu_interp(arr->cu_xb, //input data
274           sett->Ninterp, //input data length
275           arr->cu_tshift, //output time domain
276           arr->cu_xbr, //output values
277           sett->N, //output data length
278           arr->cu_d, //diagonal
279           arr->cu_dl, //upper diagonal
280           arr->cu_du, //lower diagonal
281           arr->cu_B); //coefficient matrix

```

## 2.6 Spin-down loop internals

In spin-down loop signal resampled in last step is phase-modulated by factor dependent on current spin-down parameter and shifted baricentric time.

Listing 8: `jobcore.cu`

```

325 phase_mod_2<<<(sett->N + BLOCK_SIZE - 1)/BLOCK_SIZE, BLOCK_SIZE
    >>>(
326     arr->cu_xa_f, arr->cu_xb_f,
327     arr->cu_xar_f, arr->cu_xbr_f,
328     het1, sgnlt[1], arr->cu_shft,
329     sett->N);

```

Next step is the core of whole algorithm: we compute two integrals used in evaluating  $\mathcal{F}$ -statistic and at the same time perform interpolation to get twice (or more) denser sampling rate. Interpolation can be done with two methods:



1. *interbinning*, which consists in computing values in between the bins by approximation using following formula<sup>1</sup>:  

$$x_{i+1/2} = (x_i - x_{i+1})/\sqrt{2}$$
2. *FFT interpolation*, i.e. padding with zeros to size `fftpad`-times larger than original array. This method requires computing twice as long FFT.

Listing 9: jobcore.cu

```

334     if (opts->fftinterp == INT) { //interpolation by interbinning
335         //pad zeros from N to nfft
336         pad_zeros<<<(sett->nfft - sett->N + BLOCK_SIZE - 1)/
            BLOCK_SIZE, BLOCK_SIZE>>>(
337             arr->cu_xa_f+sett->N, arr->cu_xb_f+sett->N, sett->nfft -
                sett->N);
338         CudaCheckError();
339
340         // FFT length nfft
341         CUFFT_EXEC_FFT(plans->plan, arr->cu_xa_f, arr->cu_xa_f,
            CUFFT_FORWARD);
342         CUFFT_EXEC_FFT(plans->plan, arr->cu_xb_f, arr->cu_xb_f,
            CUFFT_FORWARD);
343
344         // make a gap between every bin
345         interbinning_gap<<<(sett->nfft/2 + BLOCK_SIZE - 1)/
            BLOCK_SIZE, BLOCK_SIZE>>>(
346             arr->cu_xa_f, arr->cu_xb_f, arr->cu_xa2_f, arr->cu_xb2_f,
                sett->nfft);
347         CudaCheckError();
348
349         //interpolate values between every bin
350         interbinning_interp<<<((sett->nfft-2)/2 + BLOCK_SIZE - 1)/
            BLOCK_SIZE, BLOCK_SIZE>>>(
351             arr->cu_xa2_f, arr->cu_xb2_f, sett->nfft);
352         CudaCheckError();
353
354         cu_xa_final = arr->cu_xa2_f;
355         cu_xb_final = arr->cu_xb2_f;
356
357     } else { //interpolation by FFT (fftpad-times-longer FFT)
358         //pad zeros from N to nfft*fftpad
359         pad_zeros<<<(sett->nfft*sett->fftpad - sett->N + BLOCK_SIZE
            - 1)/BLOCK_SIZE, BLOCK_SIZE>>>(
360             arr->cu_xa_f+sett->N, arr->cu_xb_f+sett->N, sett->nfft*
                sett->fftpad - sett->N);
361         CudaCheckError();
362
363         //fft length fftpad*nfft
364         CUFFT_EXEC_FFT(plans->plan, arr->cu_xa_f, arr->cu_xa_f,
            CUFFT_FORWARD);
365         CUFFT_EXEC_FFT(plans->plan, arr->cu_xb_f, arr->cu_xb_f,
            CUFFT_FORWARD);
366
367         cu_xa_final = arr->cu_xa_f;
368         cu_xb_final = arr->cu_xb_f;
369     }

```

<sup>1</sup>I'm not sure about the sign, but later we use moduli of these values, so it's irrelevant.

Finally, having computed  $F_a$  and  $F_b$  integrals (as defined in [1], equation (5.5)) we can compute  $\mathcal{F}$ -statistic ([1], eq. (3.8)) and normalize it.

Listing 10: jobcore.cu

```

374     compute_Fstat<<<(sett->nmax-sett->nmin + BLOCK_SIZE - 1)/
        BLOCK_SIZE, BLOCK_SIZE>>>(
375         cu_xa_final + sett->nmin,
376         cu_xb_final + sett->nmin,
377         cu_F + sett->nmin,
378         sett->crf0,
379         sett->nmax - sett->nmin);

```

With assumption that noise present in signal is white noise, normalization can be done by simply dividing  $\mathcal{F}$ -statistic by signal's variance. Without that assumption  $\mathcal{F}$ -statistic is divided into blocks (size, for example, 4096 elements) and normalized by it's sum in every block<sup>2</sup>:

Listing 11: jobcore.cu

```

383     if (sett->sig2 < 0.0) { // when noise is not white-noise
384         FStat_gpu(cu_F+sett->nmin, sett->nmax - sett->nmin, NAV, arr
            ->cu_mu, arr->cu_mu_t);
385     } else { // when noise is white-noise
386         kernel_norm_Fstat_wn<<<(sett->nmax - sett->nmin + BLOCK_SIZE
            - 1)/BLOCK_SIZE, BLOCK_SIZE>>>(
387             cu_F + sett->nmin,
388             sett->sig2,
389             sett->nmax - sett->nmin);
390     }

```

## 2.7 Finding and saving candidates data

Next step consist in finding  $\mathcal{F}$ -statistic values that exceed given threshold and save found parameters in output file.

Searching is performed in parallel: every thread checks if value lying under its index is greater than neighboring values. If so it increments buffer counter and saves parameters to local buffer (i.e. spin-down-loop-internal buffer).

Listing 12: jobcore.cu

```

392     //zero candidates count
393     cudaMemset(arr->cu_cand_count, 0, sizeof(int));
394
395     /* FINDING AND SAVING CANDIDATES */
396     find_candidates<<<(sett->nmax - sett->nmin + BLOCK_SIZE - 1)/
        BLOCK_SIZE, BLOCK_SIZE>>>(
397         cu_F, // F-statistic array
398         arr->cu_cand_params, // local parameters buffer
399         arr->cu_cand_count, // local found candidates count
400         opts->trl, // threshold value
401         sett->nmin, // minimum F-statistic array index

```

<sup>2</sup>Important: 'block' length NAV described in this paragraph isn't the same as 'block' on the GPU grid length BLOCK\_SIZE.

```

402         sett->nmax,           // maximum --||--
403         sett->fftpad,         // FFT padding for calculation of freq.
404         sett->nfft,           // FFT length
405         sgnl0,                // base frequency parameter
406         sett->nd,             // number of degrees of freedom
407         sgnlt[1],             // spindown
408         sgnlt[2],             // sky\
409         sgnlt[3]);           // positions

```

Listing 13: kernels.cu

```

204 // This may be confusing, but three "if"s in kernel below
205 // are faster than one "if" with one, big, combined condition.
206 // Operations in every "if" are the same, so they are wrapped in macro.
207
208 // parameters are:
209 // [frequency, spindown, position1, position2, snr]
210 #define ADD_PARAMS_MACRO \
211     int p = atomicAdd(&found, 1); \
212     params[p*NPAR + 0] = 2.0*M_PI*(idx)/((double)fftpad*nfft)+sgnl0; \
213     params[p*NPAR + 1] = sgnl1; \
214     params[p*NPAR + 2] = sgnl2; \
215     params[p*NPAR + 3] = sgnl3; \
216     params[p*NPAR + 4] = sqrt(2*(F[idx]-ndf));
217
218
219 __global__ void find_candidates(...) {
220
221     int idx = blockIdx.x * blockDim.x + threadIdx.x + nmin;
222
223     if (idx > nmin && idx < nmax && F[idx] >= val && F[idx] > F[idx+1] &&
224         F[idx] > F[idx-1]) {
225         ADD_PARAMS_MACRO
226     } else if (idx == nmin && F[idx] >= val && F[idx] > F[idx+1]) {
227         ADD_PARAMS_MACRO
228     } else if (idx == nmax-1 && F[idx] >= val && F[idx] > F[idx-1]) {
229         ADD_PARAMS_MACRO
230     }
231 }

```

Finally found candidates parameters are copied to global buffer. If size left in global buffer is less than local buffer size (i.e. next searching could cause overflow) then we save data to output file. Saving is also performed in between first and second sky positions loop as there may not be enough found candidates to flush buffer to disk.

Listing 14: jobcore.cu

```

414 if (cand_count > 0) { //if something was found
415     printf ("\nSome signals found in %d_%d_%d_%d\n", pm, mm, nn,
416         ss);
417     cudaMemcpy(arr->cu_cand_buffer + *cand_buffer_count * NPAR,
418         arr->cu_cand_params,
419         sizeof(FLOAT_TYPE) * cand_count*NPAR,
420         cudaMemcpyDeviceToDevice);
421     *cand_buffer_count += cand_count;
422
423     // check if it's time to save results
424     // (place left in buffer is less than one-spindown buffer
425         length)

```

```

424         if (*cand_buffer_count >= arr->cand_buffer_size - arr->
           cand_params_size) {
425             save_candidates(arr->cu_cand_buffer, arr->cand_buffer,
                           cand_buffer_count, outname);
426         }
427     }

```

### 3 References

- [1] P. Astone, K. Borkowski, P. Jaranowski, M. Pietka, and A. Królak. “Data analysis of gravitational-wave signals from spinning neutron stars. V. A narrow-band all-sky search”. In: *Physical Review D* 82, 022005 (2010).